Together Serverless Image Generation Latency Benchmark

Date: March 18, 2026

Scope

This report summarizes a rough latency benchmark for five image-generation models served via Together serverless endpoints. The purpose was to estimate the time required to generate one image per model under matched input conditions.

Methodology

Benchmark runner: project image suite added to `benchmarker.py`.

Provider/endpoint: Together serverless image generation API.

Workflow tested: text-to-image only.

Shared prompt: `A premium product photo of a ceramic coffee mug on a wooden table, soft natural window light, clean background, high detail.`

Shared settings: 1024x1024 resolution, n=1, response_format=url, default model settings, and no explicit steps parameter.

Validation process: 1-run smoke test per model followed by 3 sequential runs per model for the rough benchmark.

Metric reported: wall-clock latency summarized as average, p50, and p95.

Latency Results

Rank	Model	Together model ID	Avg (s)	p50 (s)	p95 (s)	Success
1	Qwen Image 2.0	Qwen/Qwen-Image-2.0	7.53	7.19	9.27	3/3
2	Seedream 4.0	ByteDance-Seed/Seedream-4.0	8.40	8.08	9.01	3/3
3	Ideogram 3.0	ideogram/ideogram-3.0	8.67	8.69	8.71	3/3
4	FLUX.2 Pro	black-forest-labs/FLUX.2-pro	11.88	11.39	12.87	3/3
5	GPT Image 1.5	openai/gpt-image-1.5	21.55	22.20	24.47	3/3

Interpretation

For this prompt and default-serving configuration, Qwen Image 2.0 was the fastest model, followed by Seedream 4.0 and Ideogram 3.0. FLUX.2 Pro was meaningfully slower, and GPT Image 1.5 was the slowest by a wide margin.

These outputs were broadly comparable in quality, but not tightly quality-normalized. The benchmark is therefore best used as a rough latency comparison rather than a strict speed-for-equivalent-quality ranking.

Benchmark Contact Sheet

FLUX.2 Pro

Qwen Image 2.0

GPT Image 1.5

Seedream 4.0

Ideogram 3.0